

## Assignment

Your tasks are as follows. All algorithms should be your own code. No Tensorflow/Pytorch.

1. (2 pts) Apply the normalization on the training and test data.

```
) def normalize(data):
    # get mean and std.
    mean = np.mean(data, axis=0)
    std_dev = np.std(data, axis=0)

    # reenforce the division not 0
    std_dev[std_dev == 0] = 1

    # x = (x-mean) / std
    data = (data - mean) / std_dev
    return data

def preprocess_data(data):
    X = [] # (60000, 784)
    Y = [] # (60000,)
    for image_tensor, label in data:
        # Reshape each image to a flat vector of 784 elements and convert to a numpy array
        X.append(image_tensor.numpy().reshape(784))

        # rewrite the label to (0 if label between 0-4) or (1 if 5-9)
        if label <= 4:
            Y.append(0)
        else:
            Y.append(1)

    X = np.array(X, dtype=np.float32) # without enforce the dtype, may unable to process
    Y = np.array(Y, dtype=np.int64)
    X = normalize(X) # normalize feature
```

✓  
11s



```
X_train,Y_train = preprocess_data(train_data)
X_test, Y_test = preprocess_data(test_data)
print("X_train:", X_train.shape)
print("y_train:",Y_train.shape)
print("X_test:", X_test.shape)
print("Y_test:", Y_test.shape)

print(np.std(X_train))
```



```
X_train: (60000, 785)
y_train: (60000,)
X_test: (10000, 785)
Y_test: (10000,)
0.9563967519997209
```

2. (2 pts) As a baseline, train a linear classifier  $\hat{y} = v^T x$  and quadratic loss. Report its test accuracy.

```
▶ # How to implement Linear Regression from scratch with Python
class LinearRegression:
    def __init__(self, lr=0.01, epochs=10):
        self.lr = lr
        self.epochs = epochs
        self.weights = None

    def train(self, X, y):
        n_samples, n_features = X.shape
        self.weights = np.zeros(n_features)

        for epoch in range(self.epochs):
            y_pred = np.dot(X, self.weights) # f(x) = wx as linear

            loss = y - y_pred # compute loss

            dw = np.dot(X.T, loss) / n_samples
            self.weights -= self.lr * dw

            if epoch % 10 == 0: # track losses as mean square error
                print(epoch, " \niter, Loss: ", np.mean(loss**2))
                print(" iter, accuracy: ", self.accuracy(X, y))

    def predict(self, X):
        return (np.dot(X, self.weights) > 0.5) #Return if predict

    def accuracy(self, X, y):
        y_pred = self.predict(X)
        return np.mean(y_pred == y)
```



#problem 1, linear classification

```
linear_model = LinearRegression(lr=0.001, epochs=100)
```

```
linear_model.train(X_train,Y_train)
```

```
#linear_model.train(X0,Y0)
```

```
print("Accuracy:", linear_model.accuracy(X_test, Y_test))
```



0

iter, Loss: 0.49006666666666665

iter, accuracy: 0.5099333333333333

10

iter, Loss: 0.540939285159809

iter, accuracy: 0.5099333333333333

20

iter, Loss: 0.6215125870061101

iter, accuracy: 0.5066833333333334

30

iter, Loss: 0.7564380981161849

iter, accuracy: 0.44506666666666667

40

iter, Loss: 0.9938405528110311

iter, accuracy: 0.38526666666666665

50

iter, Loss: 1.4296833821470212

iter, accuracy: 0.3573

60

iter, Loss: 2.258927872976632

iter, accuracy: 0.34826666666666667

70

iter, Loss: 3.884077457454416

iter, accuracy: 0.34806666666666667

80

iter, Loss: 7.147460769903205

Observation:

I have trained with learning rate from 0.1 to 0.00001. The smaller learning rate helps reduce the loss and increase accuracy. However, it has the highest accuracy at 51%. Which kind make sense, since linear regression is not designed for classification problem, especially we labeled the target with numerical 0 and 1, model may interpret it as mathematical values.

3. (7 pts) Train a neural network classifier with quadratic loss  $\ell(y, f(x)) = (y - f(x))^2$ . Plot the progress

of the test and training accuracy (y-axis) as a function of the iteration counter  $t$  (x-axis).

Report the

final test accuracy for the following choices

- $k=5$
- $k=50$
- $k=200$
- Comment on the role of hidden units  $k$  on the ease of optimization and accuracy

```
# this structure of this neural network browsed https://github.com/OriYarden/Binary-
class NeuralNetwork:
    def __init__(self, n_neurons=10, lr=0.0001, epochs=10, loss_fun="quadratic", batch
        self.n_neurons = n_neurons
        self.lr = lr
        self.epochs = epochs

        self.loss_fun = loss_fun

        self.W = None
        self.v = None
        self.batch = batch

    # Xavier initialization for neural network
    # https://machinelearningmastery.com/weight-initialization-for-deep-learning-neura
    @staticmethod
    def xavier_initialization(n_in, n_out):
        limit = np.sqrt(1 / (n_in + n_out))
        return np.random.uniform(-limit, limit, size=(n_in, n_out))

    def initialize_weights(self, n_features):
        # for the layers, input 785 feature with 60k datapoints, after layer1 become 5,5
        np.random.seed(42)
        self.W = self.xavier_initialization(self.n_neurons, n_features)
        # self.v = self.xavier_initialization(self.n_neurons, 1)
        self.v = self.xavier_initialization(self.n_neurons, 1).reshape(-1) # make it vec
        # print(self.W.shape, "wshape")

def train(self, X, y):
    n_samples, n_features = X.shape #get number of sam
    # print(n_samples.shape)
    # print(n_features.shape)
    self.initialize_weights(n_features) #initialize w

    # print(self.initialize_weights.shape)
    # test = np.dot(X, self.W.T)
    # print(test, "forward propa")
    # print(test.shape, "forward propa")

    batch_size = self.batch
    accuracies = []

    for epoch in range(self.epochs):
        for i in range(0, X.shape[0], batch_size):
            # choose the batch datapoints
            X_batch = X[i:i + batch_size]
            y_batch = y[i:i + batch_size]

            #forward pass
            h1_input, h1_output, h1_pred = self.forward_p

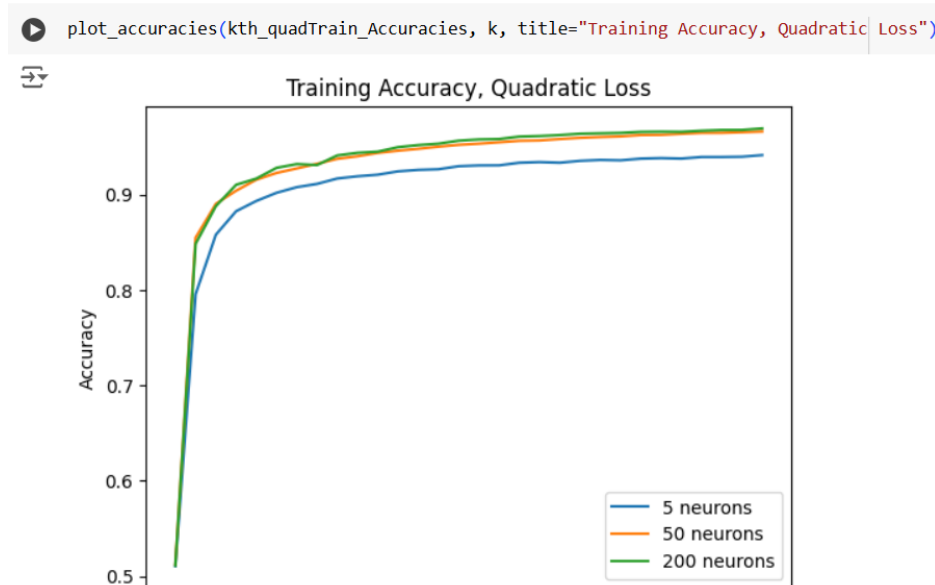
            # compute loss as quadratic fun
            if self.loss_fun == "quadratic":
                loss = y_batch - h1_pred
                # print(loss.shape)

            h1_loss = -loss # compute the h1 lavers' l
```

```
k = [5,50,200]
kth_quadTrain_Accuracies = []
for neurons in k:
    quadraticNN = NeuralNetwork(n_neurons=neurons, lr=0.001, epochs=10, batch=10, loss_fun="quadratic")
    train_accuracies = quadraticNN.train(X_train,Y_train)
    kth_quadTrain_Accuracies.append(train_accuracies)

    final_test_accuracy = quadraticNN.accuracy(X_test, Y_test)
    print("Quadratic nn for # neuron = ", neurons, " final test Accuracy: ", final_test_accuracy)

Quadratic nn for # neuron = 5 final test Accuracy: 0.9413
Quadratic nn for # neuron = 50 final test Accuracy: 0.9595
Quadratic nn for # neuron = 200 final test Accuracy: 0.9629
```



Obs:

With a common epoch =10, batch\_size=10, learning rate = 0.001

For quadratic loss function, the accuracy increases as the number of neurons increases.

Quadratic nn for # neuron = 5 final test Accuracy: 0.9413

Quadratic nn for # neuron = 50 final test Accuracy: 0.9595

Quadratic nn for # neuron = 200 final test Accuracy: 0.9629

The total number of neurons at 50 and 200 shares a very similar accuracy curve. 200 neurons doing slightly better. However, with 200 neurons the computation duration is a lot longer.

With a learning rate = 0.01, the accuracy stuck at 50%

4. (7 pts) Train a neural network classifier with logistic loss, namely  $\ell(y, f(x)) = -y \log(\sigma(f(x))) - (1 - y) \log(1 - \sigma(f(x)))$  where  $\sigma(x) = 1/(1 + e^{-x})$  is the sigmoid function. In this case, the hard-thresholding is applied on top of the sigmoid function, i.e.,  $\hat{y} = 1 \sigma(f(x)) > 0.5$ . Repeat step 3.

```

y_batch = y[1:1 + batch_size]

#forward pass
h1_input, h1_output, h1_pred = self.forward_pass(X_batch)

# compute loss as quadratic fun
if self.loss_fun == "quadratic":
    loss = y_batch - h1_pred
    # print(loss.shape)

    h1_loss = -loss # compute the h1 layers' loss as quadratic

# compute loss as logistic fun
elif self.loss_fun == "logistic":
    # use sigmoid to compute the probability first
    h1_pred_sigmoid = self.sigmoid(h1_pred)
    # then compute the loss based on some sigmoid with y_batch target
    loss = y_batch - h1_pred_sigmoid
    # print(loss.shape)

    # compute the h1 layers' loss as sigmoid
    # print("logistic before: -loss, h1Pred_Sigmoid, sigmoidDerivative")
    h1_loss = -loss * h1_pred_sigmoid * (1-h1_pred_sigmoid) #

# logistic: h1_loss, -loss, h1Pred_Sigmoid, sigmoidDerivative (1
# print("logistic after: h1_loss, -loss, h1Pred_Sigmoid, sigmoidDerivative")

# Backprop for v
dv = np.dot(h1_loss, h1_output) / batch_size
# print(self.v.shape)

```

```

k = [5,50,200]
kth_logTrain_Accuracies = []
for neurons in k:
    logisticNN = NeuralNetwork(n_neurons=neurons, lr=0.001, epochs=10, batch=10, loss_fun="logistic")
    train_accuracies = logisticNN.train(X_train,Y_train)
    kth_logTrain_Accuracies.append(train_accuracies)

# Print final test accuracy
final_test_accuracy = logisticNN.accuracy(X_test, Y_test)
print("Logistic nn for # neuron = ", neurons, " final test Accuracy: ", final_test_accuracy)

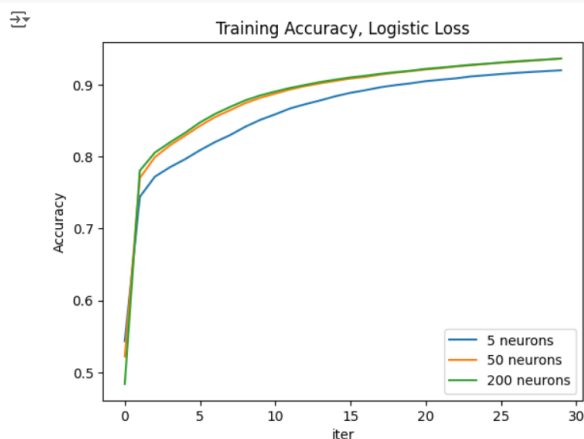
```

Logistic nn for # neuron = 5 final test Accuracy: 0.9226  
 Logistic nn for # neuron = 50 final test Accuracy: 0.9377  
 Logistic nn for # neuron = 200 final test Accuracy: 0.9372

```

plot_accuracies(kth_logTrain_Accuracies, k, title="Training Accuracy, Logistic Loss")

```



For logistic loss function, the accuracy increase as number of neurons increases.

Logistic nn for # neuron = 5 final test Accuracy: 0.9226

Logistic nn for # neuron = 50 final test Accuracy: 0.9377

Logistic nn for # neuron = 200 final test Accuracy: 0.9372

For the logistic neural network, the total number of neuron at 50 and 200 shares a very similar accuracy curve. 200 neurons doing slightly better. However, with 200 neurons the computation duration is a lot longer.

5. (2 pts) Comment on the difference between linear model and neural net. Comment on the differences between logistic and quadratic loss in terms of optimization and test/train accuracy.

The simple linear classifier has significant lower accuracy compared neural network, however, the computation is very robust. The reason Neural network overcome the linear classifier could be the increasing in complexity. The extra neurons not only helps contribute weight, they also can help contribute weights to strange images. At worst, they may not efficiently contribute the accuracy, with gradient update their corresponding weight to nearly 0 so that they would not hurt the accuracy badly.

The Logistics' accuracy is lower than the Quadratic one, and by observing the plot, at mid of the total epoch, number of neuron at 50 and 200 has no significant difference. I suspect the cause is logistic loss is more complex than quadratic loss because of sigmoid function. By choosing a smaller learning rate = 0.0001, the result is slightly better. I believe it may hit its



## limitation

```
✓ sim ▶ k = [5,50,200]
kth_logTrain_Accuracies2 = []
for neurons in k:
    nn_model = NeuralNetwork(n_neurons=neurons, lr=0.0001, epochs=100, batch=10, l
    train_accuracies = nn_model.train(X_train,Y_train)
    kth_logTrain_Accuracies2.append(train_accuracies)

    # Print final test accuracy
    final_test_accuracy = nn_model.accuracy(X_test, Y_test)
    print("Logistic nn for # neuron = ", neurons, " final test Accuracy: ", final_

# # Plot training accuracy over epochs
# plot_accuracies(train_accuracies, title="Neural Network Training Accuracy - Quad
```

```
⇒ Logistic nn for # neuron = 5 final test Accuracy: 0.923
Logistic nn for # neuron = 50 final test Accuracy: 0.9376
Logistic nn for # neuron = 200 final test Accuracy: 0.9378
```

```
✓ ls ▶ plot_accuracies(kth_logTrain_Accuracies2, k, title="Training Accuracy, Logistic Lo
```

